

OPERATING SYSTEMS, DV1697/DV1698

LAB 1: PROCESSES, THREADS, SYNCHRONIZATION, AND MEMORY SYSTEMS

Sadi Alawadi
Blekinge Institute of Technology

2024 fall
updated: 2024-10-02

This laboratory assignment is divided into two parts, and the objective of each part is to:

1. study how user-level processes and threads work, as well as study various synchronization and communication mechanisms for processes and threads.
2. how the memory system impacts the performance, with a focus on the virtual memory system and page replacement.

The laboratory assignments should be conducted, solved, implemented, and presented in groups of up to three students!

It is ok to do them individually, but groups of more than three students are not allowed.

Home Assignment 1. Plagiarism and collaboration

You are encouraged to work in groups of three. Groups larger than three are not accepted.

Discussions between laboratory groups are positive and can be fruitful. It is normally not a problem, but watch out so that you do not cross the border to cheating. For example, you are *not allowed* to share solution approaches, solutions to the different tasks, source code, output data, results, etc.

The submitted solution(s) to the laboratory assignments and tasks, should be developed by the group members only. You are not allowed to copy code from somewhere else, i.e., neither from your student fellows nor from the internet or any other source. The only other source code that is allowed to use is the one provided with the laboratory assignment.

End of home assignment 1.

1 Introduction

A computer system comprises one or more processors, with most modern processors being multicore, meaning they contain multiple CPU cores. Another critical component of any computer system is its memory system, which includes both physical main memory (a.k.a. primary memory) and secondary memory for long-term storage, as well as persistent storage (usually on disk). When a program executes, its code and data must reside in the main memory for the processor to access them. However, since main memory is limited and cannot store all the code and data for every program simultaneously, secondary memory is used to store parts of the program when needed. This process is managed through a technique called virtual memory.

In this laboratory assignment, we will explore how to divide the workload of a program using processes or threads. Additionally, we will study the memory model for processes and threads, focusing on how they share data and synchronize. We will also examine factors within the memory system that influence performance and the workings of virtual memory.

The programs for this lab assignment are tested on **Ubuntu Linux** and are designed to work in the **lab room (G332)**. However, they should also function on any Linux distribution you may have at home. If you prefer to run the lab on Windows or macOS, you can install a virtual machine, such as VirtualBox, and run Linux/Ubuntu on it. On Windows, WSL 2 (Windows Subsystem for Linux) should work for most of the tasks (although I have not tested all of them).

The source code files and programs required for the laboratory can be found on the Canvas page.

If nothing else is said/written, you shall compile your programs using:

```
gcc -O2 -o output_file_name sourcefile1.c ... sourcefileN.c
```

When you compile and link a parallel program written using pthreads, you shall add “-lpthread” to the compile command in order to link the pthreads library to your application:

```
gcc -O2 -o output_file_name sourcefile1.c ... sourcefileN.c -lpthread
```

All commands are written in a Linux terminal or console window.

2 Examination and grading

Present and discuss your solutions orally with a teacher or lab assistant. Once all tasks are completed:

- Prepare a compressed file (either in **.tar** or **.zip** format), including the source code for the working solutions of both **Part 1** and **Part 2** and submit it for evaluation. Take into account that:
 1. **Part 1** should contain the solutions of the following tasks **Task 7**, **Task 9**, **Task 11**, **Task 13**, and **Task 14**.
 2. **Part 2** should contain the solutions of the following tasks **Task 17**, **Task 20**, and **Task 23**.
- Write and submit a short report (**approximately 2-3 pages, in PDF format**) detailing your answers to the questions in the assignment and describing your implementations. Be sure to include responses to all questions posed in the tasks and any measurements, results, and relevant details.

All material (except the code given to you in this assignment) must be produced by the laboratory group alone.

Note: If further oral clarification regarding your code or report is needed, the examiner may contact you within a week. All group members must be present for the oral discussion if this occurs.

3 Part 1

Home Assignment 2. Preparations

Read through these laboratory instructions and do the **home assignments**. The purpose of the home assignments is to do them *before* the lab sessions to prepare your work and make more efficient use of the lab sessions.

Read the following sections in the course book [2]:

- Section 2.1: Processes
- Section 2.2: Threads
- Section 2.3: Interprocess communication
- Section 6.1 and 6.2: Resources and Introduction to deadlocks
- Section 8.1: Multiprocessors
- Section 10.3: Processes in Linux

End of home assignment 2.

3.1 Processes

In this first part of laboratory 1, we will study how user-level processes work. A process is created when we start a program, and processes can also create new processes.

Home Assignment 3. `fork()` and `exec()`

Study the man pages of `fork()` and `exec()` so you understand how these system calls work and their differences. (hint, e.g., write `man fork` in the terminal on a Linux system)

End of home assignment 3.

The standard way of creating new user-level processes on Unix systems is the system call `fork()`. On Linux we also have a system call `clone()`, that also creates new processes. One of the main differences between `fork()` and `clone()` is that `clone()` allows child processes to share parts of its execution context with its parent.

The program in Listing 1 is a small example of how we can create new processes using the system call `fork()`.

Task 1. Fork example

Compile and execute the program in `fork.c` (see Listing 1).

Which process is the parent and which is the child process?

The variable `i` is used in both processes. Is the value of `i` in one process affected when the other process increments it? Why/Why not?

Modify the program in two ways:

- Let the program create a third process, that writes 100 "C" (similarly to "A" and "B")
- In the *parent process*, print out the process id (pid) of each of the child processes. It is *not* allowed to use `getpid()` to get the process id of the child processes.

Which are the process identities (pid) of the child processes?

End of task 1.

3.2 Process communication and synchronization

An important property of processes is memory protection and isolation, i.e., what happens inside on process should not affect the memory and data of other processes. However, in some situations we need to pass data from one process to another, i.e., using inter process communication (IPC). One approach is *sending messages* and another is communication through a *shared memory* area.

When using shared memory for process communication, we need to allocate a piece of memory that both processes have access to. The system call to allocate such shared memory from the operating system is `shmget()`. Then each of the processes can attach or detach themselves to/from that piece of memory using `shmat()` and `shmdt()`. Shared memory communication between processes can be done using either memory-mapped files or an allocated memory segment. In our examples we will use an allocated memory segment.

Home Assignment 4. Shared memory functions

Study the man pages of `shmget()`, `shmop()`, `shmat()`, `shmdt()`, and `shmctl()` so you understand how these system calls work and their differences.

Study the program example in Listing 2 (`shmem.c`).

What does the program do?

End of home assignment 4.

In the case with only *one number* in the buffer, we have a case similar to the one described in Tanenbaum's book on pages 123–124 ('Strict Alternation'). When we increase the buffer size, we do not want to restrict the ordering between the processes to strictly alternating turns. In contrast, we would like the processes to be able to access the buffer in any order. However, we must ensure two important restrictions:

1. we cannot put a new number into the buffer if the buffer is full, and
2. we cannot remove any number if the buffer is empty.

The problem that may arise otherwise is that we may lose/drop numbers or fetch/read the same number multiple times.

Task 2. Shared memory buffer

Modify the program in Listing 2 (`shmem.c`) so the buffer contains 10 numbers instead of only one number. Implement it as a *circular, bounded buffer*, i.e.,

- the items shall be fetched by the consumer in the same order as they were put into the buffer by the producer,
- when you come to the end of the buffer you start over from the first buffer place, and
- it is not allowed to fill the buffer with 10 items before the consumer runs / fetches items in the buffer.

The producer, i.e., the parent process, shall wait a random time between 0.1s - 0.5s each time it has put a number into the buffer. Similarly, the consumer, i.e., the child process, shall also wait a random time between 0.2s - 2s each time it has fetched a number from the buffer.

Comment your program and explain where and how the problems described above can occur.

End of task 2.

In the previous example, there is a risk that both processes may read and update the same data simultaneously. Therefore, we need to have some mechanism to protect the data. A common way to do that is to introduce critical sections, where only one process at the time is allowed to access the shared data.

In the following example (`semaphore.c`), we introduce POSIX semaphores as a way to implement critical sections. POSIX semaphores come in two versions: *named semaphores*, that can be used between processes, and *unnamed semaphores*, that are memory-based and can only be used between threads in the same process or on a memory region shared between processes (such as the one that create in the previous example). In this example we will look at named semaphores.

Home Assignment 5. Semaphores

Read the man pages for `sem_open()`, `sem_post()`, `sem_wait()`, `sem_close()` and `sem_unlink()`.

Study the program example in Listing 3 (`semaphore.c`).

What does the program do?

Compile and execute the program example in Listing 3 (`semaphore.c`).

You compile it with:

```
gcc -O2 -o semaphore semaphore.c -lrt -lpthread
```

Be sure that you clearly understand how the semaphores are allocated and initialized, and how they are used.

What are the initial values of the semaphores `sem_id1` and `sem_id2`, respectively?

End of home assignment 5.

Task 3. Shared memory buffer with semaphores

Copy the program you developed in Task 2, the one with a circular buffer containing 10 slots.

Modify the new program by adding proper synchronizations using semaphores.

Make sure the program executes correctly by using semaphores, i.e., there should be no risk for dropping numbers or reading numbers several times.

Why did the problems in Task 2 occur, and how does your solution with semaphores solve them?

Hint: Think about which initial values that are suitable for the semaphores, e.g., see the example in Section 2.3.5 in the course book [2].

End of task 3.

3.3 Message queues

Sending messages is an alternative to using shared memory for inter-process communication (IPC). In the following examples we will create a process that send messages to another process using the messaging primitives available in Linux/Unix. In this type of IPC, the processes do not share memory, instead they let the kernel handle the transfer of messages between the processes. There are two different message queue APIs in Linux: System V and POSIX, and they provide similar functionality. System V message queues have been around for a long time, thus they are widely used and in essence all Unix systems implement them. POSIX message queues are newer, thus they are developed based on the experience on System V message queues.

In our laboratory, we are going to use System V message queues. The basic functionality we are interested in is the ability to send a message from one process to another and the ability to receive a message from another process. The first is implemented in a call named `msgsnd()` and the other in a call named `msgrcv()`. In this scenario, one process sends a message to the kernel, and then the kernel locates the other process and delivers the message to that process.

In the following example (`msgqsend.c` and `msgqrecv.c`), we introduce how inter-process communication can be done using System V message queues. The example is taken and modified from Tutorialspoint.¹

Home Assignment 6. Message queue example

Read the man pages for `msgget()`, `msgsnd()`, `msgrcv()`, and `msgctl()`.

Study the program example in Listing 4 (`msgqsend.c`) and Listing 5 (`msgqrecv.c`).

¹https://www.tutorialspoint.com/inter_process_communication/inter_process_communication_message_queues.htm

What do the programs do?

Compile the programs in Listing 4 (`msgqsend.c`) and Listing 5 (`msgqrecv.c`).

When you execute them, run them in two different terminal windows. Further, make sure you start `msgqsend` first.

Why do you need to start `msgqsend` first?

End of home assignment 6.

Task 4. Message queue modification

Modify the programs in Listing 4 (`msgqsend.c`) and Listing 5 (`msgqrecv.c`) to send and receive a sequence of integers instead.

A sequence of 50 integers should be generated automatically with random values in the range $0 < x < \text{INT_MAX}$.

Hint: The size of an `int` is different from the size of a `char`. You also need to think about how to detect the end of transmission, i.e., when all integers are sent and received.

Note: It is not allowed to send the integers as a string of chars. The integers should be sent as integers, one at the time (i.e., it not allowed to put them all in an array and then send the array).

End of task 4.

3.4 Threads

Threads are different to processes in how they share resources such as memory and files. A process can contain one or several threads. For example, a sequential program is executed in one process with only one single executing thread. In this laboratory we will use POSIX Threads, or `pthread`s for short.

In the first thread example, see Listing 6 (`pthreadcreate.c`), we will study how threads are created and terminated.

Home Assignment 7. Thread creation

Read the manual pages for the following functions:

- `pthread_create()`
- `pthread_exit()`
- `pthread_join()`

Especially, make sure you understand the different arguments (input parameters) to the different functions.

End of home assignment 7.

Task 5. `pthread_create()` example

Compile and execute the program in Listing 6 (`pthreadcreate.c`). Remember the `-lpthread` flag to the compiler.

Make sure you understand what it does and why.

What does the program print when you execute it?

End of task 5.

Sometimes we would like to send arguments (input parameters) to a thread we create, and sometimes we also want the thread to return some data. A tricky thing is that a thread can only have *one* argument and that argument should be of type `void*`.

In the example in Listing 7 (`pthreadcreate2.c`), we show an example of how we can pass arguments to a thread when we create the thread. Two data items, `id` and `numThreads` are put in a `struct`, and then a `void*` pointer is set to point to that `struct`, and finally, the `void*` pointer is passed to the thread function. In the child threads, the input parameter is typecasted back to the `struct` type, and then we can use the different values passed to the child thread.

Task 6. `pthread_create()` example 2

Compile and execute the program in Listing 7 (`pthreadcreate2.c`).

Make sure you understand what it does and why.

Why do we need to create a new `struct threadArgs` for each thread we create?

End of task 6.

When we want to return values from a child thread, we cannot simply execute a `return` statement as in a normal function, since the thread function must have return type `void`. A solution is to return values to the parent thread through the same `struct` as the input parameters were passed. In this situation, we cannot allow the child to deallocate the `struct` where the arguments are passed to the child thread, as in Task 6 (Listing 7). Instead, the parent thread is responsible for both allocating and deallocating the argument `struct`.

Task 7. `pthread_create()`, how to handle return values

In Listing 8 (`pthreadcreate3.c`), we provide a program where the parent thread allocates (and deallocates) an array of `structs` of type `threadArgs`, one `struct` for each child thread. The `threadArgs` `struct` passed to each child thread should also be used to pass values from the child thread to the parent thread.

Your task is to:

- Add a new variable in the `struct threadArgs`, called `squaredId`.
- In the child thread, take the thread id, square it, and return the value. Use the `threadArgs` `struct` for communicating values between the parent and the child.
- In the main program, wait for all child threads to terminate, and then print the squared id for each thread (this value should be sent to the parent thread by each child thread).

End of task 7.

3.5 Thread communication and synchronization

Since all threads in the same process share memory, they can easily communicate with each other by reading and writing data to variables in the shared memory. However, as in the case for processes, we need to protect the shared variables from simultaneous modification by different threads. A mechanism to achieve this protection is mutex variables. Therefore, we will in this part of the laboratory study how mutex synchronization works.

Home Assignment 8. Thread synchronization

Read the manual pages for the following functions:

- `pthread_mutex_init()`
- `pthread_mutex_lock()`
- `pthread_mutex_unlock()`

End of home assignment 8.

Task 8. Bank account example

In Listing 9 (`bankaccount.c`) you find a simple simulation of a bank account where deposits and withdrawals are done. Each thread does either 1000 deposits (odd thread id \Rightarrow calls `deposit()`) or 1000 withdrawals (even thread id \Rightarrow calls `withdraw()`).

If everything goes well, the saldo at the end of the execution should be 0, given an even number of threads.

Execute the program a number of times (with an even number of threads, e.g., 16, 32, 64, and 128 threads).

Does the program execute correctly? Why/why not?

End of task 8.

The problem that occurs is that two (or more) threads try to read and update the same variable at the same time, and their accesses to the shared variable (in this case `bankAccountBalance`) are interleaved. When updating a variable, a thread reads the value of the variable from the memory, updates the value (a local operation in a processor register), and finally writes the new value back to memory. When two threads do this concurrently, the execution may look like the one in Figure 1 and is called a *race condition*. Note that this situation may occur also on single-core processors, e.g., if an interrupt and process context switch happens between line 1 and 2. Programs with data races may have unpredictable behavior, and can cause bugs that are very hard to find and reproduce (so called "Hiesenbugs").

Time	Thread 1 (on cpu1)	Thread 2 (on cpu2)
1	<code>proc_reg = mem_var</code> (<code>reg_value == 1</code>)	
2		<code>proc_reg = mem_var</code> (<code>reg_value == 1</code>)
3	<code>proc_reg++</code> (<code>reg_value == 2</code>)	<code>proc_reg++</code> (<code>reg_value == 2</code>)
4	<code>mem_var = proc_reg</code> (<code>mem_value == 2</code>)	
5		<code>mem_var = proc_reg</code> (<code>mem_value == 2</code>)

Figure 1: Possible execution when two threads try to update the same variable `mem_var` concurrently, assuming that `mem_var==1` before the execution. The situation is referred to as a *race condition*. The correct value of the code sequence should be 3, not 2.

The solution to the problem in Figure 1 is to enclose accesses to shared variable in critical sections. Then we can guarantee mutual exclusion, i.e., only one thread at the time can access and modify a shared variable.

Task 9. Bank account, correction

Correct the program in Listing 9 (`bankaccount.c`) by introducing proper synchronizations on *shared variables*.

The shared variable(s) that are protected and accessed in critical sections should be *locked as short time as possible*, i.e., you are supposed to keep the critical sections as short as possible.

End of task 9.

3.6 Dining professors

The Dining Philosophers is a well-known synchronization problem, that involves potential livelock and deadlock issues. The problem is described on pages 167–170 in Tanenbaum's book [2]. In the problem, five philosophers are seated around a table and would like to eat. Each of them has a bowl and there are five chopsticks, see Figure 2. Unfortunately, they need two chopsticks in order to eat. Thus, they need to share and synchronize with each other.




Figure 2: The Dining Philosophers

In this section, we are going to implement Dining Professors. Five operating system professors, Tanenbaum, Bos, Lamport, Stallings, and Silberschatz, are seated around a round table and eat noodles. To the left of each bowl, there is a chopstick. Each professor needs two chopsticks to eat. Thus, the professors need to share the chopsticks with each other, e.g., Bos needs to share with Tanenbaum and Lamport.

Each professor sits and think for a while, and then becomes hungry. He will then pick up his left chopstick (if the chopstick is not taken by the colleague to the left). Then, the professor thinks a while before he picks up the chopstick to the right (if it is not taken). When the professor has both chopsticks, he can eat. When he has eaten for a while, he puts down both chopsticks and starts to think again.

Task 10. Dining Professors, implementation

Implement Dining Professors with five professors according to the description above. The program shall contain five threads, where each thread represents one professor. The program must guarantee that only one professor at the time can use a chopstick.

- Let the professors think a random time between one and five seconds.
- Take the left chopstick.
- When they have gotten the left chopstick, let them think in two to eight seconds.
- Take the right chopstick.
- Let the professor eat for a random time between five-ten seconds, and then put down both chopsticks.
- Write information to the terminal when the professors go from one state to another, e.g., "Tanenbaum: thinking -> got left chopstick". Also indicate when a professor tries to take a chopstick.

There are some restrictions for your implementation.

- Each fork should be modeled as a shared resource, which exclusive usage.
- You are not allowed to look at the state of your neighbors to see if they are eating or not, i.e., the solution in Section 2.5.1 in the course book [2] is not allowed.

- You are only allowed to see if the forks to the left and to the right of you are available.
- You are not allowed to lock the whole table when you examine the availability of the forks.

An implementation done in line with the description above is not deadlock-free, i.e., it may result in a deadlock. Explain why, i.e., which conditions lead to the deadlock?

End of task 10.

Task 11. Dining Professors, deadlock-free implementation

Modify the program from Task 10 so the program is deadlock-free, i.e., you can guarantee that no deadlock can occur.

Which are the four conditions for deadlock to occur?

Which of the four conditions did you remove in order to get a deadlock-free program, and how did you do that in your program?

End of task 11.

3.7 Parallelism and performance

Threads and processes are often used to introduce concurrency, i.e., the ability to do multiple things at the same time. There can be many reasons for that, e.g., handling independent requests to a server or handling different tasks such as spell checking and auto backup at the same time in a word processing application. However, another important reason to introduce concurrency is performance. When a concurrent program execute on multiple cores, it is said to execute in parallel. Ideally, if a parallel program execute on two cores it could run twice as fast, i.e., we can have a speedup of 2. Speedup is defined as:

$$\text{Speedup} = \frac{T_{\text{sequential}}}{T_{\text{parallel}}} \quad (1)$$

where $T_{\text{sequential}}$ is the execution time of sequential program and T_{parallel} is the execution time of parallel program.

Unfortunately, it is often very hard to get as high speedup as we have cores running the parallel program. When we have as high speedup as we have cores, i.e., speedup of 4 on four cores, we call it *linear speedup*. However, in most cases the speedup is sub-linear, i.e., less than the ideal one. There are many reasons for that, e.g., sequential parts of the application, synchronization overhead, scheduling overhead, etc. In the following sections we will study two different applications, *matrix multiplication* and *Quicksort*. The first one is relatively easy to get close to linear speedup, while the second one is much harder.

When measuring the execution time, it looks a bit different if you have `bash` or `tcsh` as default shell in your terminal / console. In the lab room (G332) `tcsh` is default, but in a standard Ubuntu/Linux distribution `bash` is default. You usually can find out which shell you run by executing the command (if the environment variable `SHELL` is set correctly, which is unfortunately not always done):

```
echo $SHELL
```

If you are running `bash`, you should measure the execution time with:

```
time command
```

The value of "real" is the wall clock time that we are interested in.

If you are running `tcsh`, you should measure the execution time with:

```
/usr/bin/time command
```

The value of "elapsed" is the wall clock time that we are interested in.

3.8 Parallelization of a simple application

We start by parallelizing a simple application, *matrix multiplication*. A sequential version of the program is found in the file `matmulseq.c`, see Listing 10.

Task 12. Sequential matrix multiplication

Compile and execute the program `matmulseq.c`.

How long time did it take to execute the program?

End of task 12.

Task 13. Parallel matrix multiplication

We will now parallelize the matrix multiplication program using threads. Parallelize the program according to the following assumptions:

- Parallelize only the matrix multiplication, but not the initialization.
- A new thread shall be created for each row to be calculated in the matrix.
- The row number of the row that a thread shall calculate, shall be passed to the new thread as a parameter to the new thread.
- The main function shall wait until all threads have terminated before the program terminates.

Compile and link your parallel version of the matrix multiplication program. Don't forget to add `-lpthread` when you link your program. We will now measure how much faster the parallel program is as compared to the sequential one.

Execute the program and measure the execution time.

How long was the execution time of your parallel program?

Which speedup did you get (as compared to the execution time of the sequential version, where $Speedup = T_{sequential}/T_{parallel}$)?

End of task 13.

A performance limiting factor when writing and executing parallel applications is if there is any sequential execution path in the parallel program. For example, critical sections introduce sequential parts in a program, although they are necessary for the correctness of the program. In our matrix multiplication example, we have a sequential part in the initialization of the matrices.

Task 14. Parallelization of the initialization

Parallelize the initialization of the matrices also, i.e., the function `init_matrix`. Use one thread to initialize each of the rows in the matrices `a` and `b`. Compile, link, and execute your program.

Which is the execution time and speedup for the application now?

Did the program run faster or slower now, and why?

End of task 14.

4 Part 2

Home Assignment 9. Preparations

Read through these laboratory instructions and do the **home assignments**. The purpose of the home assignments is to do them *before* the lab sessions to prepare your work and make more efficient use of the lab sessions.

Read the following sections in the course book [2]:

- Section 3.2: A memory abstraction: Address spaces
- Section 3.3: Virtual memory
- Section 3.4: Page replacement algorithms
- Section 3.5: Design issues for paging systems
- Section 3.6: Implementation issues
- Section 10.4: Memory management in Linux

End of home assignment 9.

4.1 CPU-bound vs. I/O-bound processes

In this part, we will study the differences between CPU-bound and I/O-bound processes and some aspects of the memory system behaviour.

Home Assignment 10. Test programs

Study the small test programs in Listing 11 (`test1.c`) and Listing 12 (`test2.c`), so you understand what they are doing.

How much dynamic memory is allocated in the `test1` program (i.e., how large is `struct link`)?

How much data is written to the temporary file `/tmp/file1.txt` in each iteration of the program `test2`?

End of home assignment 10.

In order to understand the performance of the virtual memory system and the disk I/O system, we need to have some way of measuring what is going on in the operating system. To our help we will use two programs, i.e., `vmstat` and `top`.

- `vmstat` provides information about, e.g., swapping activity, allocated memory, and blocks read and written to a block device.

- `top` provides information about, e.g., how much execution time, cpu utilization, and memory each process is using.

Home Assignment 11. `vmstat` and `top`

Study the man pages of `vmstat` and `top` so you understand how these two commands work. Specifically, understand the parameters that you can give to the commands and the output format of the commands.

In which output columns of `vmstat` can you find the amount of memory a process uses, the number of swap ins and outs, and the number of I/O blocks read and written?

Where in the output from `top` do you find how much cpu time and cpu utilization a process is using?

End of home assignment 11.

Task 15. CPU-bound vs. I/O-bound processes

Compile the test programs `test1` and `test2`.

Execute "`top`" and "`vmstat 1`" in two new terminal windows.

Execute the test program `test1`. How much memory does the program use? Does it correspond to your answer in Home Assignment 10?

What is the cpu utilization when executing the `test1` program? Which type of program is `test1`?

Execute the test program `test2`. How much memory does the program use? How many blocks are written out by the program? How does it correspond to your answer in Home Assignment 10?

What is the cpu utilization when executing the `test2` program? Which type of program can we consider `test2` to be?

End of task 15.

When measuring the execution time, it looks a bit different if you have `bash` or `tcsh` as default shell in your terminal / console. In the lab room (G332) `tcsh` is default, but in a standard Ubuntu/Linux distribution `bash` is default. You usually can find out which shell you run by executing the command (if the environment variable `SHELL` is set correctly, which is unfortunately not always done):

```
echo $SHELL
```

If you are running `bash`, you should measure the execution time with:

```
time command
```

The value of "real" is the wall clock time that we are interested in.

If you are running `tcsh`, you should measure the execution time with:

```
/usr/bin/time command
```

The value of "elapsed" is the wall clock time that we are interested in.

Task 16. Overlapping CPU and I/O execution

Study the two scripts `run1` (Listing 13) and `run2` (Listing 14).

What is the difference between them in terms of how they execute the test programs?

Execute the script `run1` and measure the execution time. Study the cpu utilization using `top` during the execution. How long time did it take to execute the script and how did the cpu utilization vary?

Execute the script `run2` and measure the execution time. Study the cpu utilization using `top` during the execution. How long time did it take to execute the script and how did the cpu utilization vary?

Which of the two cases executed fastest?

In both cases, the same amount of work was done. In which case was the system best utilized and why?

End of task 16.

4.2 Virtual memory and page replacement algorithms

In this part of the laboratory, we are going to study the effects of different page replacement algorithms.

Among the files provided to you, there are two files with the extension `.mem` (`mp3d.mem` and `mult.mem`). These two files contains 100,000 memory references each, where each row in the file corresponds to a memory address that the program has generated. The traces are captured when running two different programs (`mp3d`, a particle-based wind tunnel simulation, and matrix multiplication). Both applications are from a classic benchmark suite of numerical applications.

Task 17. Implementation of FIFO page replacement

Your task is to write a program that calculates the number of page faults for a sequence of memory references (i.e., the memory reference trace in the `.mem`-files) when using the FIFO (First-In-First-Out) page replacement policy.

Your program shall take as input the number of physical pages, the page size, and the name of the trace file, see the example below:

```
./fifo no_phys_pages page_size filename
```

The program shall write the resulting number of page faults for that specific combination of number of pages and page size, for either `mp3d.mem` or `mult.mem`.

Example execution:

```
mycomputer$ ./fifo 4 256 mp3d.mem
No physical pages = 4, page size = 256
Reading memory trace from mp3d.mem... Read 100000 memory references
Result: 11940 page faults
```

End of task 17.

Task 18. Page fault measurements for the FIFO page replacement policy

Use the program that you developed in Task 17 to fill in the number of pages faults in Table 1 and Table 2. To your help, we have filled in some of the numbers in the tables (so you can use them to check that you program generate the correct number of page faults).

End of task 18.

Table 1: Number of page faults for `mp3d.mem` when using FIFO as page replacement policy.

Page size	Number of pages							
	1	2	4	8	16	32	64	128
128								
256			11940					
512							417	
1024								

Table 2: Number of page faults for `mult.mem` when using FIFO as page replacement policy.

Page size	Number of pages							
	1	2	4	8	16	32	64	128
128						249		
256								
512								
1024		16855						

Task 19. Evaluation of the FIFO page replacement policy

Based on the values in Table 1 and Table 2, answer the following questions.

What is happening when we keep the number of pages constant and increase the page size? Explain why!

What is happening when we keep the page size constant and increase the number of pages? Explain why!

If we double the page size and halve the number of pages, the number of page faults sometimes decrease and sometimes increase. What can be the reason for that?

Focus now on the results in Table 2 (matmul). At some point decreases the number page faults very drastically. What memory size does that correspond to? Why does the number of page faults decrease so drastically at that point?

At some point the number of page faults does not decrease anymore when we increase the number of pages. When and why do you think that happens?

End of task 19.

Task 20. Implementation of LRU page replacement

Your task is to write a program that calculates the number of page faults for a sequence of memory references (i.e., the memory reference trace in the .mem-files) when using the LRU (Least-Recently-Used) page replacement policy.

Your program shall take as input the number of physical pages, the page size, and the name of the trace file, see the example below:

```
./lru no_phys_pages page_size filename
```

The program shall write the resulting number of page faults for that specific combination of number of pages and page size, for either mp3d.mem or mult.mem.

Example execution:

```
mycomputer$ ./lru 4 256 mp3d.mem
No physical pages = 4, page size = 256
Reading memory trace from mp3d.mem... Read 100000 memory references
Result: 9218 page faults
```

End of task 20.

Task 21. Page fault measurements for the LRU page replacement policy

Use the program that you developed in Task 20 to fill in the number of pages faults in Table 3. To your help, we have filled in some of the numbers in the table (so you can use them to check that you program generate the correct number of page faults).

End of task 21.

Task 22. Comparison of the FIFO and LRU page replacement policies

Based on the values in Table 1 and Table 3, answer the following questions.

Table 3: Number of page faults for mp3d.mem when using LRU as replacement policy.

Page size	Number of pages							
	1	2	4	8	16	32	64	128
128								
256			9218					
512							362	
1024								

Which of the page replacement policies FIFO and LRU seems to give the lowest number of page faults? Explain why!

In some of the cases, the number of page faults are the same for both FIFO and LRU. Which are these cases? Why is the number of page faults equal for FIFO and LRU in those cases? Explain why!

End of task 22.

Task 23. Implementation of Optimal page replacement (Bélády's algorithm)

Your task is to write a program that calculates the minimum number of page faults for a sequence of memory references, i.e., you should implement the Optimal page replacement policy (which is also known as Bélády's algorithm).

Your program shall take as input the number of physical pages, the page size, and the name of the trace file, see the example below:

```
./optimal no_phys_pages page_size filename
```

The program shall write the resulting number of page faults for that specific combination of number of pages and page size, for either mp3d.mem or mult.mem.

Example execution:

```
mycomputer$ ./optimal 32 128 mp3d.mem
No physical pages = 32, page size = 128
Reading memory trace from mp3d.mem... Read 100000 memory references
Result: 824 page faults
```

End of task 23.

Task 24. Page fault measurements for the Optimal (Bélády's) page replacement policy

Use the program that you developed in Task 23 to fill in the number of pages faults in Table 4. To your help, we have filled in some of the numbers in the table (so you can use them to check that you program generate the correct number of page faults).

End of task 24.

Table 4: Number of page faults for mp3d.mem when using Optimal (Bélády's algorithm) as replacement policy.

Page size	Number of pages							
	1	2	4	8	16	32	64	128
128						824		
256								
512								
1024			3367					

Task 25. Comparison of the FIFO, LRU, and Optimal page replacement policies

Based on the values in Table 1, Table 3, and Table 4, answer the following questions.

As expected, the Optimal policy gives the lowest number of page faults. Explain why!

Optimal is considered to be impossible to use in practice. Explain why!

Does FIFO and/or LRU have the same number of page faults as Optimal for some combination(s) of page size and number of pages? If so, for which combination(s) and why?

End of task 25.

References

- [1] Andrew S. Tanenbaum and Herbert Bos, "Modern Operating Systems, 4th ed", *Pearson Education Limited*, ISBN-10: 0-13-359162-X, 2015.
- [2] Andrew S. Tanenbaum and Herbert Bos, "Modern Operating Systems, 5th Global Edition", *Pearson Education Limited*, ISBN-10: 1292459662.

A Program listings

Listing 1. fork.c

```
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    pid_t pid;
    unsigned i;
    unsigned niterations = 100;
    pid = fork();
    if (pid == 0) {
        for (i = 0; i < niterations; ++i)
            printf("A = %d, ", i);
        printf("In child => Own pid : %d\n", getpid());
        printf("In child => Parent's pid : %d\n", pid);
    } else {
        for (i = 0; i < niterations; ++i)
            printf("B = %d, ", i);
        printf("In child => Own pid : %d\n", getpid());
        printf("In child => Parent's pid : %d\n", pid);
    }
    printf("\n");

    // }
    // else{
    // printf("In Parent => Child's pid is %d\n", pid);
}
```

Listing 2. shmem.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#define SHMSIZE 128
#define SHM_R 0400
#define SHM_W 0200

int main(int argc, char **argv)
{
    struct shm_struct {
        int buffer;
        unsigned empty;
    };
    volatile struct shm_struct *shmp = NULL;
    char *addr = NULL;
    pid_t pid = -1;
    int var1 = 0, var2 = 0, shmid = -1;
    struct shm_ds *shm_buf;

    /* allocate a chunk of shared memory */
    shmid = shmget(IPC_PRIVATE, SHMSIZE, IPC_CREAT | SHM_R | SHM_W);
    shmp = (struct shm_struct *) shmat(shmid, addr, 0);
    shmp->empty = 0;
    pid = fork();
    if (pid != 0) {
        /* here's the parent, acting as producer */
        while (var1 < 100) {
            /* write to shmem */

```

```

        var1++;
        while (shmp->empty == 1); /* busy wait until the buffer is empty */
        printf("Sending %d\n", var1); fflush(stdout);
        shmp->buffer = var1;
        shmp->empty = 1;
    }
    shmdt(addr);
    shmctl(shmid, IPC_RMID, shm_buf);
} else {
    /* here's the child, acting as consumer */
    while (var2 < 100) {
        /* read from shmem */
        while (shmp->empty == 0); /* busy wait until there is something */
        var2 = shmp->buffer;
        shmp->empty = 0;
        printf("Received %d\n", var2); fflush(stdout);
    }
    shmdt(addr);
    shmctl(shmid, IPC_RMID, shm_buf);
}
}

```

Listing 3. semaphore.c

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <semaphore.h>
#include <fcntl.h> /* For O_* constants */

const char *semName1 = "my_sema1";
const char *semName2 = "my_sema2";

int main(int argc, char **argv)
{
    pid_t pid;
    sem_t *sem_id1 = sem_open(semName1, O_CREAT, O_RDWR, 1);
    sem_t *sem_id2 = sem_open(semName2, O_CREAT, O_RDWR, 0);
    int i, status;

    pid = fork();
    if (pid) {
        for (i = 0; i < 100; i++) {
            sem_wait(sem_id1);
            putchar('A'); fflush(stdout);
            sem_post(sem_id2);
        }
        sem_close(sem_id1);
        sem_close(sem_id2);
        wait(&status);
        sem_unlink(semName1);
        sem_unlink(semName2);
    } else {
        for (i = 0; i < 100; i++) {
            sem_wait(sem_id2);
            putchar('B'); fflush(stdout);
            sem_post(sem_id1);
        }
        sem_close(sem_id1);
        sem_close(sem_id2);
    }
}

```

Listing 4. msgqsend.c

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

#define PERMS 0644
struct my_msgbuf {
    long mtype;
    char mtext[200];
};

int main(void) {
    struct my_msgbuf buf;
    int msqid;
    int len;
    key_t key;
    system("touch msgq.txt");

    if ((key = ftok("msgq.txt", 'B')) == -1) {
        perror("ftok");
        exit(1);
    }

    if ((msqid = msgget(key, PERMS | IPC_CREAT)) == -1) {
        perror("msgget");
        exit(1);
    }
    printf("message queue: ready to send messages.\n");
    printf("Enter lines of text, ^D to quit:\n");
    buf.mtype = 1; /* we don't really care in this case */

    while (fgets(buf.mtext, sizeof buf.mtext, stdin) != NULL) {
        len = strlen(buf.mtext);
        /* remove newline at end, if it exists */
        if (buf.mtext[len-1] == '\n') buf.mtext[len-1] = '\0';
        if (msgsnd(msqid, &buf, len+1, 0) == -1) /* +1 for '\0' */
            perror("msgsnd");
    }
    strcpy(buf.mtext, "end");
    len = strlen(buf.mtext);
    if (msgsnd(msqid, &buf, len+1, 0) == -1) /* +1 for '\0' */
        perror("msgsnd");

    if (msgctl(msqid, IPC_RMID, NULL) == -1) {
        perror("msgctl");
        exit(1);
    }
    printf("message queue: done sending messages.\n");
    return 0;
}

```

Listing 5. msgqrecv.c

```

#include <stdio.h>
#include <stdlib.h>

```

```

#include <string.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

#define PERMS 0644
struct my_msgbuf {
    long mtype;
    char mtext[200];
};

int main(void) {
    struct my_msgbuf buf;
    int msqid;
    int toend;
    key_t key;

    if ((key = ftok("msgq.txt", 'B')) == -1) {
        perror("ftok");
        exit(1);
    }

    if ((msqid = msgget(key, PERMS)) == -1) { /* connect to the queue */
        perror("msgget");
        exit(1);
    }
    printf("message queue: ready to receive messages.\n");

    for(;;) { /* normally receiving never ends but just to make conclusion */
        /* this program ends with string of end */
        if (msgrcv(msqid, &buf, sizeof(buf.mtext), 0, 0) == -1) {
            perror("msgrcv");
            exit(1);
        }
        printf("recvd: \"%s\"\n", buf.mtext);
        toend = strcmp(buf.mtext, "end");
        if (toend == 0)
            break;
    }
    printf("message queue: done receiving messages.\n");
    system("rm msgq.txt");
    return 0;
}

```

Listing 6. pthreadcreate.c

```

#include <stdio.h>
#include <pthread.h> // pthread types and functions

void* child() {
    printf("This is the child thread.\n");
}

int main(int argc, char** argv) {
    pthread_t thread; // struct for child-thread info
    // spawn thread:
    pthread_create(&thread, // the handle for it
        NULL, // its attributes
        child, // the function it should run
        NULL); // args to that function

    printf("This is the parent (main) thread.\n");
    pthread_join(thread, NULL); // wait for child to finish
    return 0;
}

```

Listing 7. pthreadcreate2.c

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

struct threadArgs {
    unsigned int id;
    unsigned int numThreads;
};

void* child(void* params) {
    struct threadArgs *args = (struct threadArgs*) params;
    unsigned int childID = args->id;
    unsigned int numThreads = args->numThreads;
    printf("Greetings from child #%u of %u\n", childID, numThreads);
    free(args);
}

int main(int argc, char** argv) {
    pthread_t* children; // dynamic array of child threads
    struct threadArgs* args; // argument buffer
    unsigned int numThreads = 0;
    // get desired # of threads
    if (argc > 1)
        numThreads = atoi(argv[1]);
    children = malloc(numThreads * sizeof(pthread_t)); // allocate array of handles
    for (unsigned int id = 0; id < numThreads; id++) {
        // create threads
        args = malloc(sizeof(struct threadArgs));
        args->id = id;
        args->numThreads = numThreads;
        pthread_create(&(children[id]), // our handle for the child
            NULL, // attributes of the child
            child, // the function it should run
            (void*)args); // args to that function
    }
    printf("I am the parent (main) thread.\n");
    for (unsigned int id = 0; id < numThreads; id++) {
        pthread_join(children[id], NULL);
    }
    free(children); // deallocate array
    return 0;
}

```

Listing 8. pthreadcreate3.c

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

struct threadArgs {
    unsigned int id;
    unsigned int numThreads;
};

void* child(void* params) {
    struct threadArgs *args = (struct threadArgs*) params;
    unsigned int childID = args->id;

```

```

    unsigned int numThreads = args->numThreads;
    printf("Greetings from child #%u of %u\n", childID, numThreads);
}

int main(int argc, char** argv) {
    pthread_t* children; // dynamic array of child threads
    struct threadArgs* args; // argument buffer
    unsigned int numThreads = 0;
    // get desired # of threads
    if (argc > 1)
        numThreads = atoi(argv[1]);
    children = malloc(numThreads * sizeof(pthread_t)); // allocate array of handles
    args = malloc(numThreads * sizeof(struct threadArgs)); // args vector
    for (unsigned int id = 0; id < numThreads; id++) {
        // create threads
        args[id].id = id;
        args[id].numThreads = numThreads;
        pthread_create(&(children[id]), // our handle for the child
            NULL, // attributes of the child
            child, // the function it should run
            (void*)&args[id]); // args to that function
    }
    printf("I am the parent (main) thread.\n");
    for (unsigned int id = 0; id < numThreads; id++) {
        pthread_join(children[id], NULL);
    }
    free(args); // deallocate args vector
    free(children); // deallocate array
    return 0;
}

```

Listing 9. bankaccount . c

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

// Shared Variables
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
double bankAccountBalance = 0;

void deposit(double amount) {
    bankAccountBalance += amount;
}

void withdraw(double amount) {
    bankAccountBalance -= amount;
}

// utility function to identify even-odd numbers
unsigned odd(unsigned long num) {
    return num % 2;
}

// simulate id performing 1000 transactions
void do1000Transactions(unsigned long id) {
    for (int i = 0; i < 1000; i++) {
        if (odd(id))
            deposit(100.00); // odd threads deposit
        else
            withdraw(100.00); // even threads withdraw
    }
}

void* child(void* buf) {
    unsigned long childID = (unsigned long)buf;

```



```

    do1000Transactions(childID);
    return NULL;
}

int main(int argc, char** argv) {
    pthread_t *children;
    unsigned long id = 0;
    unsigned long nThreads = 0;
    if (argc > 1)
        nThreads = atoi(argv[1]);
    children = malloc( nThreads * sizeof(pthread_t) );
    for (id = 1; id < nThreads; id++)
        pthread_create(&(children[id-1]), NULL, child, (void*)id);
    do1000Transactions(0); // main thread work (id=0)
    for (id = 1; id < nThreads; id++)
        pthread_join(children[id-1], NULL);
    printf("\nThe final account balance with %lu threads is $%.2f.\n\n", nThreads, bankAccountBalance);
    free(children);
    pthread_mutex_destroy(&lock);
    return 0;
}

```

Listing 10. matmulseq.c

```

/*****
 *
 * Sequential version of Matrix–Matrix multiplication
 *
 *****/

#include <stdio.h>
#include <stdlib.h>

#define SIZE 1024

static double a[SIZE][SIZE];
static double b[SIZE][SIZE];
static double c[SIZE][SIZE];

static void
init_matrix(void)
{
    int i, j;

    for (i = 0; i < SIZE; i++)
        for (j = 0; j < SIZE; j++) {
            /* Simple initialization, which enables us to easy check
             * the correct answer. Each element in c will have the same
             * value as SIZE after the matmul operation.
             */
            a[i][j] = 1.0;
            b[i][j] = 1.0;
        }
}

static void
matmul_seq()
{
    int i, j, k;

    for (i = 0; i < SIZE; i++) {
        for (j = 0; j < SIZE; j++) {
            c[i][j] = 0.0;
            for (k = 0; k < SIZE; k++)
                c[i][j] = c[i][j] + a[i][k] * b[k][j];
        }
    }
}

```

```

    }
}

static void
print_matrix(void)
{
    int i, j;

    for (i = 0; i < SIZE; i++) {
        for (j = 0; j < SIZE; j++)
            printf(" %7.2f", c[i][j]);
        printf("\n");
    }
}

int
main(int argc, char **argv)
{
    init_matrix();
    matmul_seq();
    //print_matrix();
}

```

Listing 11. test1.c

```

#include <stdio.h>
#include <stdlib.h>
/* test1.c */

#define SIZE (32*1024)
#define ITERATIONS 10

int main(int argc, char **argv)
{
    struct link {
        int x[SIZE][SIZE];
    };
    struct link *start;
    int iteration, i, j;
    start = (struct link *) calloc(1, sizeof(struct link));
    if (!start) {
        printf("Fatal error: Can't allocate memory of %d x %d = %lu\n", SIZE, SIZE, (unsigned long)SIZE*SIZE);
        exit(-1);
    }
    for (iteration = 0; iteration < ITERATIONS; iteration++) {
        printf("test1, iteration: %d\n", iteration);
        for (i = 0; i < SIZE; i++)
            for (j = 0; j < SIZE; j++)
                start->x[i][j] = 0;
    }
}

```

Listing 12. test2.c

```

#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
/* test2.c */

```

```
#define SIZE (16*1024)
#define ITERATIONS 5

int main(int argc, char **argv)
{
    struct link {
        int x[SIZE][SIZE];
    };
    struct link *start;
    int iteration;
    FILE *f;
    start = (struct link *) calloc(1, sizeof(struct link));
    if (!start) {
        printf("Fatal error: Can't allocate memory of %d x %d = %lu\n", SIZE, SIZE, (unsigned long)SIZE*SIZE);
        exit(-1);
    }
    for (iteration = 0; iteration < ITERATIONS; iteration++) {
        printf("test2, iteration: %d\n", iteration);
        f = fopen("/tmp/file1.txt", "w");
        fwrite(start->x, sizeof(start->x[0][0]), SIZE*SIZE, f);
        fclose(f);
    }
}
```

Listing 13. run1

```
#!/bin/bash
./test1; ./test1; ./test1; ./test2; ./test2
```

Listing 14. run2

```
#!/bin/bash
(./test1; ./test1; ./test1) &
(./test2; ./test2)
```
